

Reference images + creative controls



Design-space sampler — Latin hypercube, 4 axes



Brief compiler — template + LLM enrichment



DIFFUSION SYNTHESIS

Latent diffusion synthesis — CLIP-conditioned, n stills



Quality gate — 5 calibrated checks



LEARNED REPRESENTATION

Frozen encoders — CLIP/SigLIP, DINOv2, ArcFace



SCORING

Prediction head — Bradley-Terry objective



Weighted scoring and sorting — image-stage rank



DIFFUSION SYNTHESIS

Video diffusion — promoted subset only



Delivery — platform safe areas, end card